AI Career Copilot

The program that opens today. 21 October 2026.

What it is

This is a local page on this computer. It reads one CV, names the kinds of work that experience supports, searches for jobs that are still open, and shows how each job matches. The person reviews every file and sends the application. The program does not apply.

The brain is the grok program already signed in on this computer. The page does not call https://api.x.ai and it does not use an API key. It cannot read the weekly quota, so it does not enforce a spending cap.

How to open it

From the project folder, after grok login has been done on this computer:
.\.venv\Scripts\career-copilot.exe ui
Open http://127.0.0.1:8501. Stop the command to close the page. Do not set XAI_API_KEY.

What is built

1. User

The person uploads one CV, a DOCX or a single-column text PDF. A scanned page or a multi-column file is refused in one sentence. Read the CV asks grok for up to eight roles and up to twelve skill names. A skill is kept only when that name appears in the CV. The person keeps roles, removes roles, or types one. Save writes the roles. Continue opens Search.

2. Search

No question is required. A blank field is not a preference and adds nothing to a search. Experience level, employment type, and workplace each accept more than one choice. Country, city, and languages are shown. No country means worldwide. A typed language removes an automatic result written in another language, and a posting that requires a language that was not typed. The language is read from the description after the site menu is removed. An English description that only mentions another language as a plus stays. Other notes is an optional box. Those lines guide the search and the match. They do not remove a job by themselves. Save writes the choices. Continue opens Jobs.

3. Jobs

Opening Jobs does not search. Automatic search asks grok to look on the public web. The call may use up to three turns, and it stops after two searches. The whole search stops at 4 minutes. The query is the saved roles, the saved skills, and the Search fields that are filled. The CV is not part of that search. Python sends it when it scores each opened page. It is not put into the web query, and it is not stored in the job list.

Python opens each https link again. It drops LinkedIn, private addresses, dead pages, and login walls. It then removes the site menu, including a block of short lines that repeats, and a long topic menu above the vacancy. A short list that appears once stays. Dash lines, country lists, language switchers, and consent comments are removed. The job description stays. The stored text is cut on a line at 8,000 characters. A page that says applications are closed, or whose deadline has passed, is left out. A date such as 15 Nov '26 counts, including when it is on the line after Deadline. If no deadline is found, the job stays and Notes says so. A description written in a language that was not saved is left out. A job is kept only when the posting shares at least one skill or role written in the CV. Scoring stops when three jobs reach 50 percent, or when the 4 minutes end.

The table keeps Title, Company, Found at, and Match. It shows at most three jobs at 50 percent or above. If every scored job is under 50 percent, the table shows the single highest match and, under the table: Nada foi encontrado, então retornamos a vaga com o maior match com o seu CV. If no page was opened, the line is only: Nada foi encontrado.

Opening a row shows the title, the link, and the cleaned posting, then Match, What matches, Possible matches, What is missing, Notes, and Documents. The posting shows the deadline, hours, salary, location, what the job asks, and what it offers. Yes moves a line into What matches. No moves it into What is missing. The percent is recalculated then. The same answer is used on the next search. A line with no answer is not saved. Keep this job, beside the title, keeps the job. Remove this job drops that link from a later search.

Documents always names the CV. A cover letter, portfolio, transcript, writing sample, certificate, or ID appears only when the posting names it. Generate CV for this job writes one Amsterdam PDF: A4, one column, no photo, and no

table. Experience, education, publications, languages, and contact stay word for word. Only the summary and Technical Skills may change, and only with facts already in the CV. A skill may use the job's wording when the CV already contains that skill. If the model text cannot be used, the PDF is still written and the page says that the summary and the skills were left as written in the CV. The file name opens the preview, and the page offers the PDF download. Notes beside that file name changes only the summary and Technical Skills. A note about a locked section is refused and the file stays as it was. Generate cover letter for this job appears only when a cover letter or a letter of motivation is named. It has no Notes button. The other named documents say to bring your own file. Each job keeps its own files. After the CV file exists, the page asks Did you apply? Yes or Not yet is stored with that vacancy. Nothing is submitted.

4. Documents

Documents lists every CV written on this computer. Each row shows the vacancy, the company, the date the CV was written, and whether the person said they applied. The row opens or downloads that CV, and the cover letter when one was written. Yes and Not yet change the stored answer. The list is a file on this computer. It does not store the CV text, and the program does not send the application.

Saved on this computer

roles.json stores the chosen roles. skills.json stores the skill names. search.json stores the filters and the other notes. It does not store the CV. cv.txt stores the CV text. A written job file does not change it. jobs.json stores the latest matches, the jobs the person kept, and the links that were removed. traits.json stores each possible match the person accepted or refused. Each analysis reads that file. The Jobs page does not show it. documents/ stores the CV and cover letter written for one job. applications.json stores the vacancy, the company, the date, and whether the person said they applied. It does not store the CV text.

Next search change

The Jobs page still uses one grok call, at most three turns, and two web searches. These six points are the next change. They are not in the program yet. Python will build three or four short queries from the saved roles, the saved skills, and the Search fields that are filled. Each query groups skills, role names, the filled place, and the words job, vacancy, or vacature. The CV text, an empty field, and the year stay out. One grok call runs for each query, with at most two turns and about 45 seconds. It returns only JSON with the title, the company, and an https url. It does not score the CV. Python saves the links as soon as that call returns. A later call that stops leaves those links. The page says the search took too long only when no real link was saved. Python then opens each new https page and drops a closed page, a past deadline, a removed link, a description in a language that was not saved, and a page that shares no skill or role with the CV. It scores the remaining pages one by one with the lighter call, using the CV and traits.json. It stops at three jobs at 50 percent or above, or when the 4 minutes end. A score already made stays, including the single highest job under 50 percent. The 4-minute cap, the list of three jobs, the 50 percent rule, and the sentence under the table stay as they are.

Still to do

No further writing phase is waiting. The person reviews every file and sends the application. The program does not submit it.

What it does not do

It does not submit an application. It does not invent a job, a company, or a link. It does not invent an employer, a title, a date, a number, or a skill. It does not treat a missing deadline as a closed job. It does not put a public website, accounts, or hosted storage in front of the CV.